

```
# import Automatic Differentiation
# You may use Neural Network Framework, but only for building MLPs
# i.e. no fancy probabilistic implementations
import Pkg;Pkg.add("Flux")
```

```
Resolving package versions...
  Updating `~/julia/environments/v1.3/Project.toml`
[no changes]
  Updating `~/julia/environments/v1.3/Manifest.toml`
[no changes]
```

```
import Pkg;Pkg.add("MLDatasets")
```

```
Resolving package versions...
  Updating `~/julia/environments/v1.3/Project.toml`
[no changes]
  Updating `~/julia/environments/v1.3/Manifest.toml`
[no changes]
```

```
import Pkg;Pkg.add("BSON")
```

```
Resolving package versions...
  Updating `~/julia/environments/v1.3/Project.toml`
[no changes]
  Updating `~/julia/environments/v1.3/Manifest.toml`
[no changes]
```

```
import Pkg;Pkg.add("Images")
```

```
Resolving package versions...
  Updating `~/julia/environments/v1.3/Project.toml`
[no changes]
  Updating `~/julia/environments/v1.3/Manifest.toml`
[no changes]
```

```
using Flux
using MLDatasets
using Statistics
using Logging
using Test
using Random
using StatsFuns: log1pexp
Random.seed!(412414);

#### Probability Stuff
# Make sure you test these against a standard implementation!
function skillcontour!(f; colour=nothing)
    n = 100
```

```

x = range(-5,stop=3,length=n)
y = range(-5,stop=3,length=n)
z_grid = Iterators.product(x,y) # meshgrid for contour
z_grid = reshape.(collect.(z_grid),:,1) # add single batch dim
z = f.(z_grid)
z = getindex.(z,1)
max_z = maximum(z)
levels = [.99, 0.9, 0.8, 0.7,0.6,0.5, 0.4, 0.3, 0.2] .* max_z
if colour==nothing
p1 = contour!(x, y, z, fill=false, levels=levels)
else
p1 = contour!(x, y, z, fill=false, c=colour,levels=levels,colorbar=false)
end
plot!(p1)
end

function plot_line_equal_skill!()
plot!(range(-3, 3, length=200), range(-3, 3, length=200), label="Equal Skill")
end
# log-pdf of x under Factorized or Diagonal Gaussian  $N(x|\mu,\sigma I)$ 
function factorized_gaussian_log_density(mu, logsig,xs)
"""
mu and logsig either same size as x in batch or same as whole batch
returns a 1 x batchsize array of likelihoods
"""
sigma = exp.(logsig)
return sum((-1/2)*log.(2*pi*sigma.^2) .+ -1/2 * ((xs .- mu).^2)./(sigma.^2),dim
s=1)
end

# log-pdf of x under Bernoulli
function bernoulli_log_density(logit_means,x)
"""Numerically stable log_likelihood under bernoulli by accepting  $\mu/(1-\mu)$ """
b = x .* 2 .- 1 # {0,1} -> {-1,1}
return - log1pexp.(-b .* logit_means)
end
## This is really bernoulli
@testset "test stable bernoulli" begin
using Distributions
x = rand(10,100) .> 0.5
mu = rand(10)
logit_mu = log.(mu./(1 .- mu))
@test logpdf.(Bernoulli.(mu),x) ≈ bernoulli_log_density(logit_mu,x)
# over i.i.d. batch
@test sum(logpdf.(Bernoulli.(mu),x),dims=1) ≈ sum(bernoulli_log_density(logit_mu,x),dims=1)
end

```

Test Summary:		Pass	Total
test stable bernoulli		2	2

```
# sample from Diagonal Gaussian  $x \sim N(\mu, \sigma I)$  (hint: use reparameterization
trick here)
sample_diag_gaussian( $\mu$ , log $\sigma$ ) = ( $\epsilon$  = randn(size( $\mu$ ));  $\mu$  .+ exp.(log $\sigma$ ).* $\epsilon$ )
# sample from Bernoulli (this can just be supplied by library)
sample_bernoulli( $\theta$ ) = rand.(Bernoulli.( $\theta$ ))

# Load MNIST data, binarise it, split into train and test sets (10000 e
ach) and partition train into mini-batches of  $M=100$ .
# You may use the utilities from A2, or dataloaders provided by a frame
work
function load_binarized_mnist(train_size, test_size)
    train_x, train_label = MNIST.traindata(1:train_size);
    test_x, test_label = MNIST.testdata(1:test_size);
    @info "Loaded MNIST digits with dimensionality $(size(train_x))"
    train_x = reshape(train_x, 28*28,:)
    test_x = reshape(test_x, 28*28,:)
    @info "Reshaped MNIST digits to vectors, dimensionality $(size(train_
x))"
    train_x = train_x .> 0.5; #binarize
    test_x = test_x .> 0.5; #binarize
    @info "Binarized the pixels"
    return (train_x, train_label), (test_x, test_label)
end

function batch_data((x,label)::Tuple, batch_size=100)
    """
    Shuffle both data and image and put into batches
    """
    N = size(x)[end] # number of examples in set
    rand_idx = shuffle(1:N) # randomly shuffle batch elements
    batch_idx = Iterators.partition(rand_idx,batch_size) # split into bat
ches
    batch_x = [x[:,i] for i in batch_idx]
    batch_label = [label[i] for i in batch_idx]
    return zip(batch_x, batch_label)
end
# if you only want to batch xs
batch_x(x::AbstractArray; batch_size=100) = first.(batch_data((x,zeros(
size(x)[end])),batch_size))
function skillcontour!(f; colour=nothing)
    n = 100
    x = range(-3,stop=3,length=n)
    y = range(-3,stop=3,length=n)
    z_grid = Iterators.product(x,y) # meshgrid for contour
    z_grid = reshape.(collect.(z_grid),:,1) # add single batch dim
    z = f.(z_grid)
    z = getindex.(z,1)'
    max_z = maximum(z)
```

```

levels = [.99, 0.9, 0.8, 0.7, 0.6, 0.5, 0.4, 0.3, 0.2] .* max_z
if colour==nothing
p1 = contour!(x, y, z, fill=false, levels=levels)
else
p1 = contour!(x, y, z, fill=false, c=colour, levels=levels, colorbar=false)
end
plot!(p1)
end

function plot_line_equal_skill!()
plot!(range(-3, 3, length=200), range(-3, 3, length=200), label="Equal Skill")
end

### Implementing the model

## Load the Data
batch_size = 10000
train_data, test_data = load_binarized_mnist(batch_size, batch_size)
train_x, train_label = train_data;
test_x, test_label = test_data;

## Test the dimensions of loaded data
@testset "correct dimensions" begin
@test size(train_x) == (784, batch_size)
@test size(train_label) == (batch_size,)
@test size(test_x) == (784, batch_size)
@test size(test_label) == (batch_size,)
end

```

Test Summary:		Pass	Total
correct dimensions		4	4

```

## Model Dimensionality
# ##### Set up model according to Appendix C (using Bernoulli decoder for Binarized MNIST)
# Set latent dimensionality=2 and number of hidden units=500.
Dz, Dh = 2, 500
Ddata = 28^2

# ## Generative Model
# This will require implementing a simple MLP neural network
# See example_flux_model.jl for inspiration
# Further, you should read the Basics section of the Flux.jl documentation
# https://fluxml.ai/Flux.jl/stable/models/basics/
# that goes over the simple functions you will use.
# You will see that there's nothing magical going on inside these neural network libraries
# and when you implemented a neural network in previous assignments you

```

```

    did most of the work.
# If you want more information about how to use the functions from Flux
, you can always reference
# the internal docs for each function by typing `?` into the REPL:
# ? Chain
# ? Dense
#Q1(b)
decoder = Chain(Dense(Dz, Dh, tanh), Dense(Dh, Ddata))
## Model Distributions
#Q1(a)
log_prior(z) = factorized_gaussian_log_density(0, 0, z)#TODO

#Q1(c)
function log_likelihood(x,z)
    """ Compute log likelihood log_p(x|z)"""
    d = decoder(z)# TODO: parameters decoded from latent z
    return sum(bernoulli_log_density(d, x), dims = 1) # return likelihood
    for each element in batch
end

#Q1(d)
joint_log_density(x,z) = log_prior(z) .+ log_likelihood(x,z)#TODO

## Amortized Inference
function unpack_gaussian_params(θ)
    μ, logσ = θ[1:2,:], θ[3:end,:]
    return μ, logσ
end
#Q2(a)
encoder = Chain(Dense(Ddata,Dh,tanh), Dense(Dh,2*Dz), unpack_gaussian_params)#TODO
# Hint: last "layer" in Chain can be 'unpack_gaussian_params'

#Q2(b)
log_q(q_μ, q_logσ, z) = factorized_gaussian_log_density(q_μ, q_logσ, z)
#TODO: write log likelihood under variational distribution.

#Q2(c)
function elbo(x)
    q_μ, q_logσ = encoder(x)#TODO variational parameters from data
    z = sample_diag_gaussian(q_μ, q_logσ)#TODO: sample from variational distribution
    log_joint_ll = joint_log_density(x,z) #TODO: log joint density of z and x under model
    log_q_z = log_q(q_μ, q_logσ, z)#TODO: log likelihood of z under variational distribution
    elbo_estimate = mean(log_joint_ll - log_q_z) #TODO: Scalar value, mean variational evidence lower bound over batch
    return elbo_estimate
end

#Q2(d)

```

```

function loss(x)
    return -elbo(x)#TODO: scalar value for the variational loss over elements in the batch
end

# Training with gradient optimization:
# See example_flux_model.jl for inspiration

#Q2(e)
function train_model_params!(loss, encoder, decoder, train_x, test_x; n_epochs=10)
    # model params
    ps = Flux.params(encoder,decoder)#TODO parameters to update with gradient descent
    # ADAM optimizer with default parameters
    opt = ADAM()
    # over batches of the data
    for i in 1:n_epochs
        # compute gradients with respect to variational loss over batch
        for d in batch_x(train_x)
            gs = Flux.gradient(ps) do
                batch_loss = loss(d)
                return batch_loss
            end
            #update the paramters with gradients
            Flux.Optimise.update!(opt,ps,gs)
        end
        if i%1 == 0 # change 1 to higher number to compute and print less frequently
            @info "Test loss at epoch $i: $(loss(batch_x(test_x)[1]))"
        end
    end
    @info "Parameters of encoder and decoder trained!"
end

## Load the trained model!
using BSON:@load
cd(@__DIR__)
@info "Changed directory to $(@__DIR__)"
load_dir = "trained_models"
@load joinpath(load_dir,"encoder_params.bson") encoder
@load joinpath(load_dir,"decoder_params.bson") decoder
@info "Load model params from $load_dir"

### Save the trained model!
using BSON:@save
cd(@__DIR__)
@info "Changed directory to $(@__DIR__)"
save_dir = "trained_models"
if !isdir(save_dir)
    mkdir(save_dir)
    @info "Created save directory $save_dir"
end

```

```

@save joinpath(save_dir,"encoder_params.bson") encoder
@save joinpath(save_dir,"decoder_params.bson") decoder
@info "Saved model params in $save_dir"

## Train the model(already done)
#train_model_params!(loss,encoder,decoder,train_x,test_x, nepochs=10)
@show "final elbo is $(loss(batch_x(test_x)[1]))"

```

```

"final elbo is $(loss((batch_x(test_x))[1]))" = "final elbo is 159.819
04829
745534"


```

```


# Visualization
using Images
using Plots
# make vector of digits into images, works on batches also
mnist_img(x) = ndims(x)==2 ? Gray.(reshape(x,28,28,:)) : Gray.(reshape(
x,28,28))

## Example for how to use mnist_img to plot digit from training data
plot(mnist_img((train_x[:,1])' [1,:]))
p = plot(mnist_img(train_x[:,1]))
#Q3(a)
plot_list1 = Any[]
plot_list2 = Any[]
for i in 1:10
    z = randn(2,)
    log_mean = decoder(z)
    ber_mean = exp.(log_mean) ./ (1 .+exp.(log_mean))
    samp = sample_bernoulli(ber_mean)
    push!(plot_list1, plot(mnist_img(ber_mean) ))
    push!(plot_list2, plot(mnist_img(samp) ))
end
plot_list = vcat(plot_list1, plot_list2)
display(plot(plot_list..., layout = grid(2, 10), size = (1000, 800)))

```



```
#Q3(b)
num = 10000
mean_vec = encoder(train_x[:,1:num])[1]
display(scatter(mean_vec[1,:], mean_vec[2,:], group = train_label[1:num]
))
```

```

#Q3(c)
function interpolate(za,zb,α)
    return α .* za .+ (1-α) .* zb
end
function find_first_occ(x,A)
    for i in 1:length(A)
        if A[i] == x
            return i
        end
    end
    return -1
end
function bern_mean(z)
    logit_mean = decoder(z)
    return exp.(logit_mean) ./ (1 .+ exp.(logit_mean))
end
function plot_latent(z)
    return plot(mnist_img(bern_mean(z)))
end
#plot interpolation graphs for label a and b
function plot_interp(a,b)
    result = Any[]
    indices = [find_first_occ(i,train_label) for i in [a,b]]
    pair_X = train_x[:,indices]
    #display(plot(mnist_img(pair_X[:,1])))
    push!(result,plot(mnist_img(pair_X[:,1])))
    means = encoder(pair_X)[1]
    α = 0:0.1:1
    for alpha in α
        z = interpolate(means[:,2], means[:,1],alpha)
    end
end


```

```

    plot = plot_latent(z)
    #display(plot)
    push!(result, plot)
end
#display(plot(mnist_img(pair_X[:,2]))')
push!(result, plot(mnist_img(pair_X[:,2]))')
return result
end
final_plot = vcat(plot_interp(4,2), plot_interp(3,9), plot_interp(0,1))
@show "I've included the original binary images as well for comparison.
"
```

"I've included the original binary images as well for comparison." = "
I've
included the original binary images as well for comparison."

```
display(plot(final_plot..., layout = grid(3,13), size = (1000,800)))
```




```

#Q4(a)
Dtop = convert{Int,0.5*Ddata}
#note: although this function accepts only 784*1 arrays, I've written a
nother function
#to help plot only half the graph
function tophalf(A)
    return A[1:Dtop,:].'
```

```

end
#function to plot half of the graph
mnist_img_half(x) = ndims(x)==2 ? Gray.(reshape(x,28,14,:)) : Gray.(res
hape(x,28,14))
a = plot(mnist_img(train_x[:,1]))'
b = plot(mnist_img_half(tophalf(train_x[:,1])[:,1] )')
display(plot([a,b]..., layout = grid(1,2), size = (500,400)))

```



```

## Example, left is original, right is tophalf
#p(xtop | z)
function logp_tophalf(xtop,z)
    d = decoder(z)
    return sum(bernoulli_log_density(d[1:Dtop,:], xtop), dims = 1)
end
function log_joint_tophalf(xtop,zs)
    return log_prior(zs) .+ logp_tophalf(xtop,zs)
end
#Q4(b)
#pick digit = 1
train_dig1_ind = [i for i in 1:num if train_label[i] == 1]
train_label1 = train_label[train_dig1_ind]
train_x1 = train_x[:, train_dig1_ind]
M = (length(train_label1) ÷ 100)*100
#initialize variational parameters
q_mean, q_logsig = param(rand(Dz,)) , param(rand(Dz,))
#compute elbo estimate
function neg_elbo_top(xtop, q_mean, q_logsig)
    #q_means = repeat(q_mean, 1, K)
    #q_logsiqs = repeat(q_logsig, 1, K)

```

```

#sum = 0
#for k in 1:K
    z = sample_diag_gaussian(q_mean, q_logsig)
    #z = sample_diag_gaussian(q_means, q_logsig)
    log_joint_ll = log_joint_tophalf(xtop, z) #TODO: log joint density
of z and x under model
    log_q_z = factorized_gaussian_log_density(q_mean, q_logsig, z)#TODO
: log likelihood of z under variational distribution
    elbo_est = mean(log_joint_ll .- log_q_z)
    #sum = sum + elbo_est
#end
return -elbo_est
end
# function to train the model
function train_tophalf_para(train_x1, q_mean, q_logsig, loss_function;n
= 10)
    # ADAM optimizer with default parameters
    ps = Flux.params([q_mean, q_logsig])
    opt = ADAM()
    # over batches of the data
    x_first = tophalf(batch_x(train_x1[:,1:M])[1])
    for i in 1:n
        for x in batch_x(train_x1[:,1:M])
            xtop = tophalf(x)
            gs = Flux.gradient(() -> loss_function(xtop), ps)
            #gs = Flux.gradient(ps) do
            # return loss_function(xtop)
            #end
            Flux.Optimise.update!(opt,ps,gs)
        end
        if i%20 == 0 # change 1 to higher number to compute and print less
frequently
            @info "Variational Test loss at epoch $i: $(loss_function(x_first
))"
        end
    end
end
end
q_mean, q_logsig = randn(2,), randn(2,)
@info "qmean, qlogsig is $((q_mean, q_logsig))"
loss_top(xtop) = neg_elbo_top(xtop, q_mean, q_logsig)
#training parameters, already done
#train_tophalf_para(train_x1,q_mean,q_logsig, loss_top; n = 1000)
@info "qmean, qlogsig is $((q_mean, q_logsig))"

#load training models
using BSON:@load
cd(@__DIR__)
@info "Changed directory to $(@__DIR__)"
load_dir = "trained_models"
@load joinpath(load_dir,"q_mean.bson") q_mean
@load joinpath(load_dir,"q_logsig.bson") q_logsig
@info "Load model params from $load_dir"

```

```

@info "Load model params from $save_dir"

### Save the trained model!
using BSON:@save
cd(@__DIR__)
@info "Changed directory to $(@__DIR__)"
save_dir = "trained_models"
if !isdir(save_dir)
    mkdir(save_dir)
    @info "Created save directory $save_dir"
end
@save joinpath(save_dir,"q_mean.bson") q_mean
@save joinpath(save_dir,"q_logsig.bson") q_logsig
@info "Saved model params in $save_dir"

Xtop = tophalf(batch_x(train_x1[:,1:M])[1])
@show "final test loss is $(loss_top(Xtop))"

```

```


"final test loss is $(loss_top(Xtop))" = "final test loss is 58.144529
91509
796"

```

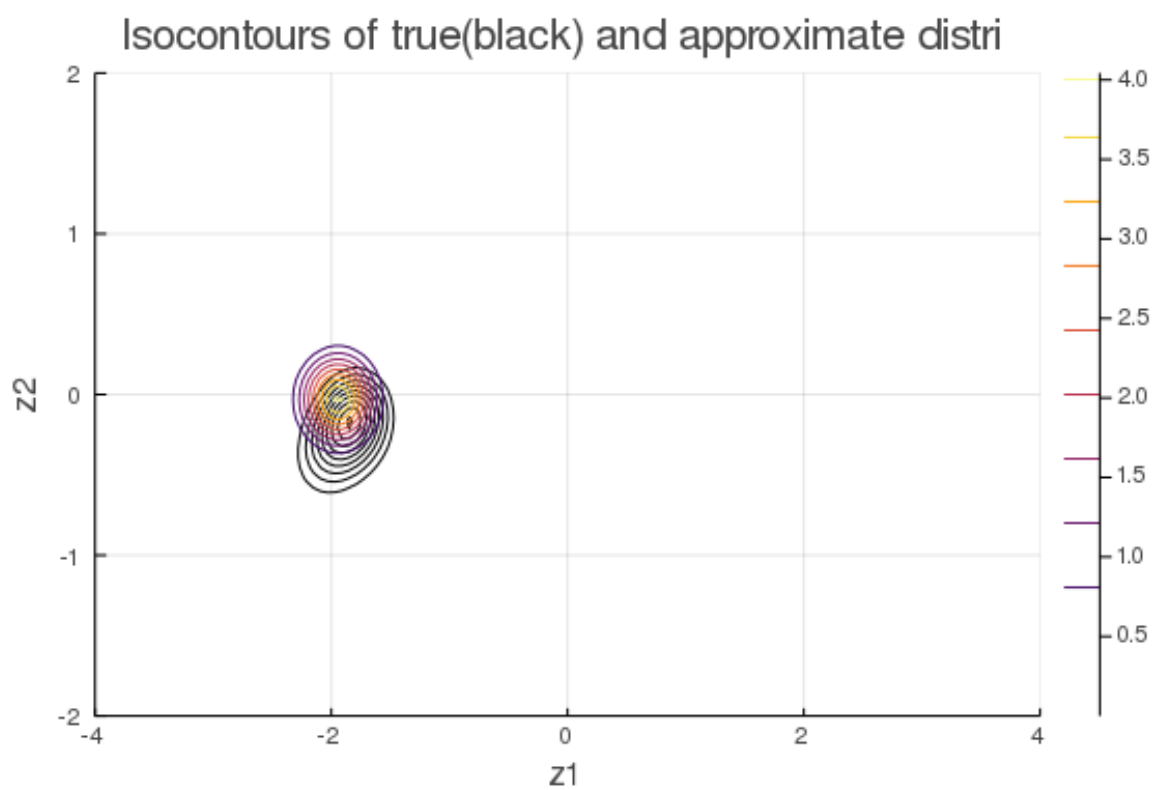
```

#sample a random z
zs = randn(Dz) .* exp.(q_logsig) .+ q_mean
logit_mean = decoder(zs)
ber_mean = exp.(logit_mean) ./ (exp.(logit_mean) .+ 1)
lower = 393:784
lowerhalf_mean = ber_mean[lower]
#plot the original image and variational one
plot_lower = plot(mnist_img_half(lowerhalf_mean[:,1] ))
original = train_x1[:,1]
plot_original = plot(mnist_img(original))
plot_upper = plot(mnist_img_half(original[1:Dtop]))
plot_appr = plot([plot_upper,plot_lower]..., layout = grid(2,1))
display(plot([plot_appr, plot_original]..., layout = grid(1,2)))

```



```
# plotting isocontours
#using all data
plot(title="Isocontours of true(black) and approximate distri",
      xlabel = "z1",
      ylabel = "z2",
      xlim = [-4,4],
      ylim = [-2,2]
)
X = tophalf(batch_x(train_x1[:,1:M])[1])
true_joint(zs) = exp.(mean(log_joint_tophalf(X, zs)))
skillcontour!(true_joint)
appr_inf(zs) = exp.(mean(factorized_gaussian_log_density(q_mean, q_logs
ig,zs )))
display(skillcontour!(appr_inf))
```



```
#Q4(c)  
print("My answer is TFFFT.")
```

My answer is TFFFT.

Published from vae.jl using Weave.jl on 2020-04-18.

